

# DNN - Experiment 04

- SIA VASHIST
- PRN: 20190802107

```
In [8]: import tensorflow as tf
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Load the iris dataset
iris = load_iris()

In [9]: # Convert the dataset to a pandas DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

In [10]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.3, random_state=42)

# Standardize the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

In [11]: # Define the DNN model
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=[4]),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(3)
])

# Compile the model
model.compile(loss='mse', optimizer='adam')

In [12]: # Train the model
model.fit(X_train, y_train, epochs=20)

# Evaluate the model on the testing set
loss = model.evaluate(X_test, y_test)
print('Test loss:', loss)

Epoch 1/20
4/4 [=====] - 2s 6ms/step - loss: 1.4979
Epoch 2/20
4/4 [=====] - 0s 7ms/step - loss: 1.2613
Epoch 3/20
4/4 [=====] - 0s 6ms/step - loss: 1.0532
Epoch 4/20
4/4 [=====] - 0s 6ms/step - loss: 0.8619
Epoch 5/20
4/4 [=====] - 0s 7ms/step - loss: 0.6768
Epoch 6/20
4/4 [=====] - 0s 8ms/step - loss: 0.5306
Epoch 7/20
4/4 [=====] - 0s 7ms/step - loss: 0.3959
Epoch 8/20
4/4 [=====] - 0s 8ms/step - loss: 0.2910
Epoch 9/20
4/4 [=====] - 0s 6ms/step - loss: 0.2177
Epoch 10/20
4/4 [=====] - 0s 6ms/step - loss: 0.1806
Epoch 11/20
4/4 [=====] - 0s 6ms/step - loss: 0.1629
Epoch 12/20
4/4 [=====] - 0s 5ms/step - loss: 0.1552
Epoch 13/20
4/4 [=====] - 0s 6ms/step - loss: 0.1486
Epoch 14/20
4/4 [=====] - 0s 7ms/step - loss: 0.1377
Epoch 15/20
4/4 [=====] - 0s 7ms/step - loss: 0.1246
Epoch 16/20
4/4 [=====] - 0s 8ms/step - loss: 0.1142
Epoch 17/20
4/4 [=====] - 0s 7ms/step - loss: 0.1043
Epoch 18/20
4/4 [=====] - 0s 7ms/step - loss: 0.0981
Epoch 19/20
4/4 [=====] - 0s 6ms/step - loss: 0.0916
Epoch 20/20
4/4 [=====] - 0s 6ms/step - loss: 0.0845
2/2 [=====] - 0s 13ms/step - loss: 0.0801
Test loss: 0.08012888580560684

In [13]: # Predict the target values using the trained model
y_pred = model.predict(X_test)

# Print the predicted and actual target values for the first 10 samples
print('Predicted targets:', np.argmax(y_pred[:10], axis=1))
print('Actual targets:', y_test[:10])

2/2 [=====] - 0s 3ms/step
Predicted targets: [1 2 2 2 2 1 2 0 2 2]
Actual targets: [1 0 2 1 1 0 1 2 1 1]
```


## Observation:

The code then defines a DNN model using TensorFlow's Keras API, with two hidden layers of 64 neurons each and a final output layer with 3 neurons corresponding to the 3 target classes in the iris dataset. The model is compiled using mean squared error (MSE) as the loss function and Adam as the optimizer.

```
In [14]: #Visualisation:  
import matplotlib.pyplot as plt  
  
history = model.fit(X_train, y_train, epochs=100, validation_data=(X_test, y_test))  
plt.plot(history.history['loss'], label='Training loss')  
plt.plot(history.history['val_loss'], label='Testing loss')  
plt.legend()  
plt.show()
```

Epoch 1/100  
4/4 [=====] - 0s 61ms/step - loss: 0.0787 - val\_loss: 0.0782  
Epoch 2/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0738 - val\_loss: 0.0747  
Epoch 3/100  
4/4 [=====] - 0s 28ms/step - loss: 0.0688 - val\_loss: 0.0715  
Epoch 4/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0644 - val\_loss: 0.0683  
Epoch 5/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0612 - val\_loss: 0.0642  
Epoch 6/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0580 - val\_loss: 0.0614  
Epoch 7/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0554 - val\_loss: 0.0593  
Epoch 8/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0533 - val\_loss: 0.0577  
Epoch 9/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0511 - val\_loss: 0.0572  
Epoch 10/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0498 - val\_loss: 0.0565  
Epoch 11/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0483 - val\_loss: 0.0550  
Epoch 12/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0469 - val\_loss: 0.0519  
Epoch 13/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0463 - val\_loss: 0.0494  
Epoch 14/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0456 - val\_loss: 0.0483  
Epoch 15/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0449 - val\_loss: 0.0478  
Epoch 16/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0436 - val\_loss: 0.0479  
Epoch 17/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0430 - val\_loss: 0.0480  
Epoch 18/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0426 - val\_loss: 0.0486  
Epoch 19/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0425 - val\_loss: 0.0479  
Epoch 20/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0416 - val\_loss: 0.0457  
Epoch 21/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0414 - val\_loss: 0.0441  
Epoch 22/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0410 - val\_loss: 0.0432  
Epoch 23/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0407 - val\_loss: 0.0427  
Epoch 24/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0407 - val\_loss: 0.0441  
Epoch 25/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0397 - val\_loss: 0.0444  
Epoch 26/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0394 - val\_loss: 0.0438  
Epoch 27/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0390 - val\_loss: 0.0442  
Epoch 28/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0387 - val\_loss: 0.0438  
Epoch 29/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0383 - val\_loss: 0.0450  
Epoch 30/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0383 - val\_loss: 0.0468  
Epoch 31/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0385 - val\_loss: 0.0460  
Epoch 32/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0381 - val\_loss: 0.0443  
Epoch 33/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0373 - val\_loss: 0.0435  
Epoch 34/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0369 - val\_loss: 0.0434  
Epoch 35/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0368 - val\_loss: 0.0426  
Epoch 36/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0363 - val\_loss: 0.0438  
Epoch 37/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0362 - val\_loss: 0.0458  
Epoch 38/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0364 - val\_loss: 0.0453  
Epoch 39/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0364 - val\_loss: 0.0444  
Epoch 40/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0360 - val\_loss: 0.0439  
Epoch 41/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0357 - val\_loss: 0.0413  
Epoch 42/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0351 - val\_loss: 0.0412  
Epoch 43/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0352 - val\_loss: 0.0419  
Epoch 44/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0347 - val\_loss: 0.0402  
Epoch 45/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0348 - val\_loss: 0.0388  
Epoch 46/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0348 - val\_loss: 0.0386  
Epoch 47/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0342 - val\_loss: 0.0399  
Epoch 48/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0348 - val\_loss: 0.0414  
Epoch 49/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0345 - val\_loss: 0.0405  
Epoch 50/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0341 - val\_loss: 0.0401  
Epoch 51/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0341 - val\_loss: 0.0394  
Epoch 52/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0337 - val\_loss: 0.0395  
Epoch 53/100  
4/4 [=====] - 0s 28ms/step - loss: 0.0340 - val\_loss: 0.0388  
Epoch 54/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0340 - val\_loss: 0.0405  
Epoch 55/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0335 - val\_loss: 0.0401  
Epoch 56/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0337 - val\_loss: 0.0393

Epoch 57/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0337 - val\_loss: 0.0396  
Epoch 58/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0333 - val\_loss: 0.0395  
Epoch 59/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0331 - val\_loss: 0.0392  
Epoch 60/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0330 - val\_loss: 0.0387  
Epoch 61/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0329 - val\_loss: 0.0374  
Epoch 62/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0329 - val\_loss: 0.0366  
Epoch 63/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0327 - val\_loss: 0.0369  
Epoch 64/100  
4/4 [=====] - 0s 30ms/step - loss: 0.0329 - val\_loss: 0.0366  
Epoch 65/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0331 - val\_loss: 0.0356  
Epoch 66/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0325 - val\_loss: 0.0368  
Epoch 67/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0329 - val\_loss: 0.0376  
Epoch 68/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0326 - val\_loss: 0.0364  
Epoch 69/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0327 - val\_loss: 0.0367  
Epoch 70/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0326 - val\_loss: 0.0376  
Epoch 71/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0326 - val\_loss: 0.0379  
Epoch 72/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0321 - val\_loss: 0.0366  
Epoch 73/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0324 - val\_loss: 0.0348  
Epoch 74/100  
4/4 [=====] - 0s 27ms/step - loss: 0.0320 - val\_loss: 0.0361  
Epoch 75/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0320 - val\_loss: 0.0381  
Epoch 76/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0319 - val\_loss: 0.0372  
Epoch 77/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0318 - val\_loss: 0.0363  
Epoch 78/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0319 - val\_loss: 0.0357  
Epoch 79/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0320 - val\_loss: 0.0351  
Epoch 80/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0317 - val\_loss: 0.0351  
Epoch 81/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0313 - val\_loss: 0.0349  
Epoch 82/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0313 - val\_loss: 0.0347  
Epoch 83/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0313 - val\_loss: 0.0354  
Epoch 84/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0313 - val\_loss: 0.0347  
Epoch 85/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0313 - val\_loss: 0.0341  
Epoch 86/100  
4/4 [=====] - 0s 26ms/step - loss: 0.0311 - val\_loss: 0.0336  
Epoch 87/100  
4/4 [=====] - 0s 25ms/step - loss: 0.0310 - val\_loss: 0.0341  
Epoch 88/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0309 - val\_loss: 0.0356  
Epoch 89/100  
4/4 [=====] - 0s 24ms/step - loss: 0.0313 - val\_loss: 0.0372  
Epoch 90/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0312 - val\_loss: 0.0366  
Epoch 91/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0307 - val\_loss: 0.0362  
Epoch 92/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0307 - val\_loss: 0.0345  
Epoch 93/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0308 - val\_loss: 0.0351  
Epoch 94/100  
4/4 [=====] - 0s 21ms/step - loss: 0.0304 - val\_loss: 0.0361  
Epoch 95/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0307 - val\_loss: 0.0359  
Epoch 96/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0307 - val\_loss: 0.0346  
Epoch 97/100  
4/4 [=====] - 0s 21ms/step - loss: 0.0303 - val\_loss: 0.0332  
Epoch 98/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0305 - val\_loss: 0.0334  
Epoch 99/100  
4/4 [=====] - 0s 22ms/step - loss: 0.0300 - val\_loss: 0.0334  
Epoch 100/100  
4/4 [=====] - 0s 23ms/step - loss: 0.0301 - val\_loss: 0.0344



## Observation:

The model is trained on the training set using the fit method and evaluated on the testing set using the evaluate method. The predicted target values are then computed using the trained model and printed alongside the actual target values for the first 10 samples.

---

## Conclusion:

The given code is an example of how to build and train a deep neural network (DNN) model for classification on the iris dataset using TensorFlow's Keras API. The code demonstrates how to prepare the dataset using scikit-learn, define a DNN model with an input layer, two hidden layers, and an output layer, and train and evaluate the model on the dataset. The code also emphasizes the importance of splitting the dataset into training and testing sets to evaluate the model's performance on unseen data.